

Lecture 8:

DBSCAN Clustering and Analytics Dataflow

**COSC 526:
Introduction to Data Mining
Spring 2021**



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE
KNOXVILLE

BIG ORANGE. BIG IDEAS.®

Lecture Outline

- Assignment 8
 - Three problems with missing data
- Project Discussion
 - Define key questions and tentative title
- Dataflow and DBSCAN
 - More in the next slides
- If time left
 - Live chat and video - https://www.youtube.com/watch?v=_2u_eHHzRto

Lecture Outline

- Use work in paper “*Development of a Scalable Method for Creating Food Groups Using the NHANES Dataset and MapReduce*” to learn about:
 - Using a dataflow for tackling a data problem
 - Structuring a research project into a set of slides
 - Motivation, goals, background, methodology and results
 - Using a different clustering method than k-mean
 - How to cluster numerical data with the Density-based spatial clustering of applications with noise (DBSCAN)
 - How to set up the setting parameters of the DBSCAN
 - Using code from other scientists
 - Re-use rather than rewriting from scratch
 - Replicability of work as the first step for new research

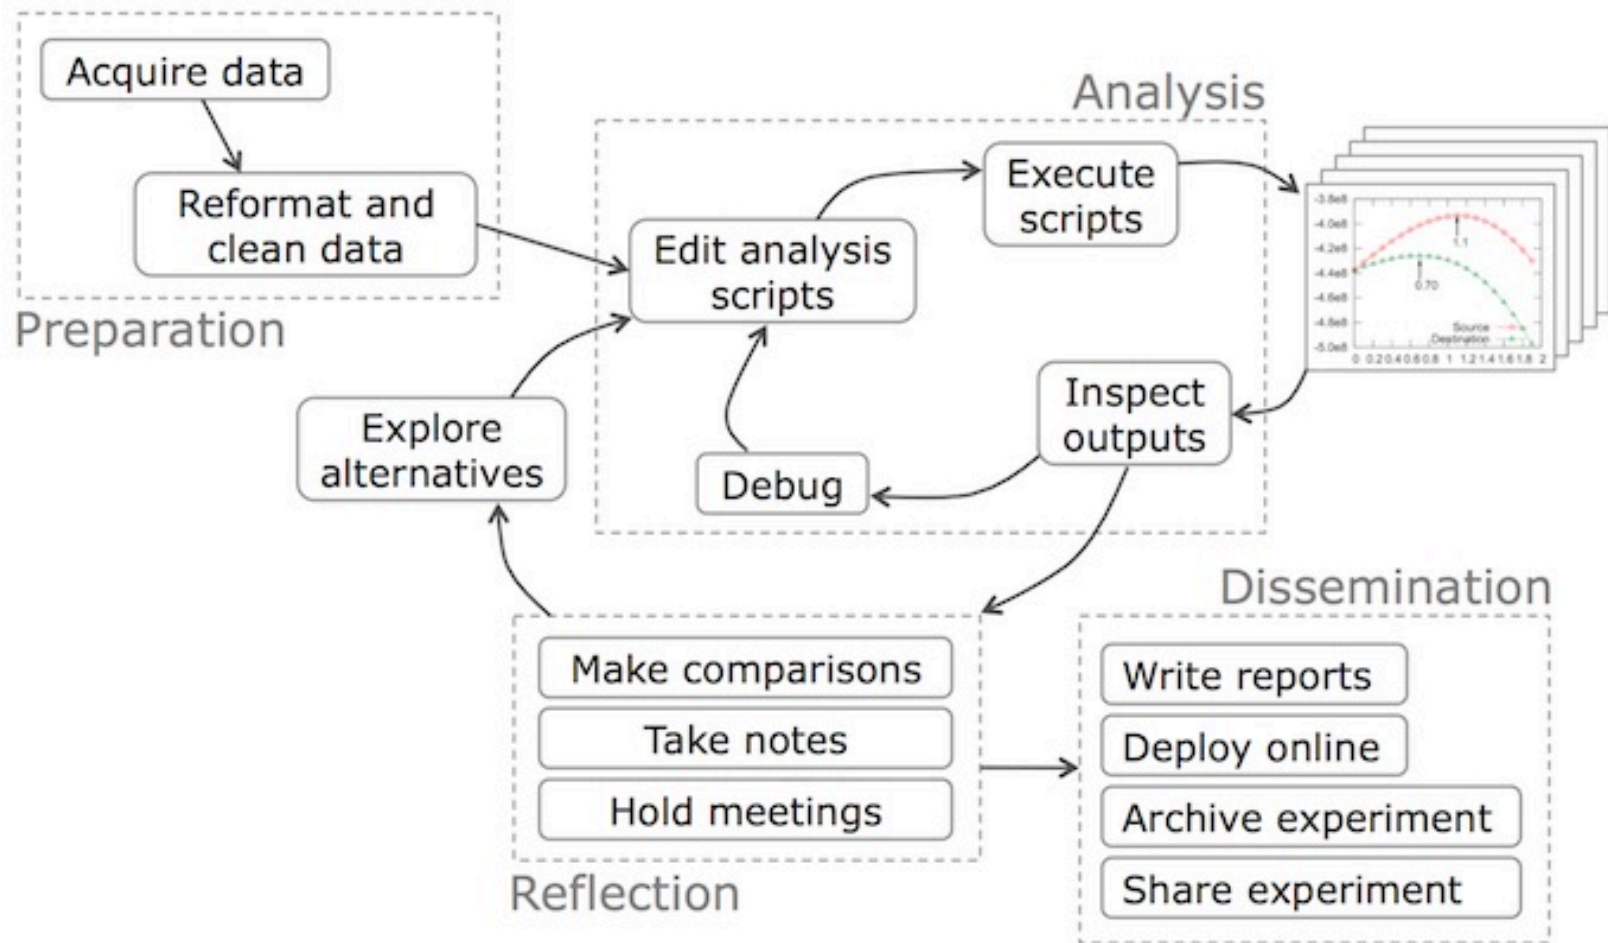
Reading

Reading

- Philip Guo. Data Science Workflow: Overview and Challenges. BLOG@CACM, October 30, 2013
- Philip Guo. Software Tools to Facilitate Research Programming. PhD Dissertation, 2012
- Michael R. Wyatt II, Travis Johnston, Mia Papas, and Michela Taufer. Development of a Scalable Method for Creating Food Groups Using the NHANES Dataset and MapReduce. In proceedings of IEEE eScience, 2019.

Data Workflow


Data Workflow



Clustering Food Items

USDA Food Classification

- Used in dietary datasets to assign food items to groups
- Subjective and general
- Categorical, not nutrient-driven



What's Wrong with USDA Food Groups?

- Food Groups are **subjective**
 - Based on human expertise
 - Based on dietary trends
- Using **subjective** food groups for dietary studies give **bad results**
- Lack of standard food groups for use in dietary studies



How Do We Fix Food Groups?

- Food Groups should be **objective**
 - Based on micro- and macro-nutrient content
- We need a standard set of food groups for dietary studies
- We need **scalable** methods for identifying food groups

“Define an objective methods to group food items in a dietary datasets based on the item micro- and macro-nutrient content”

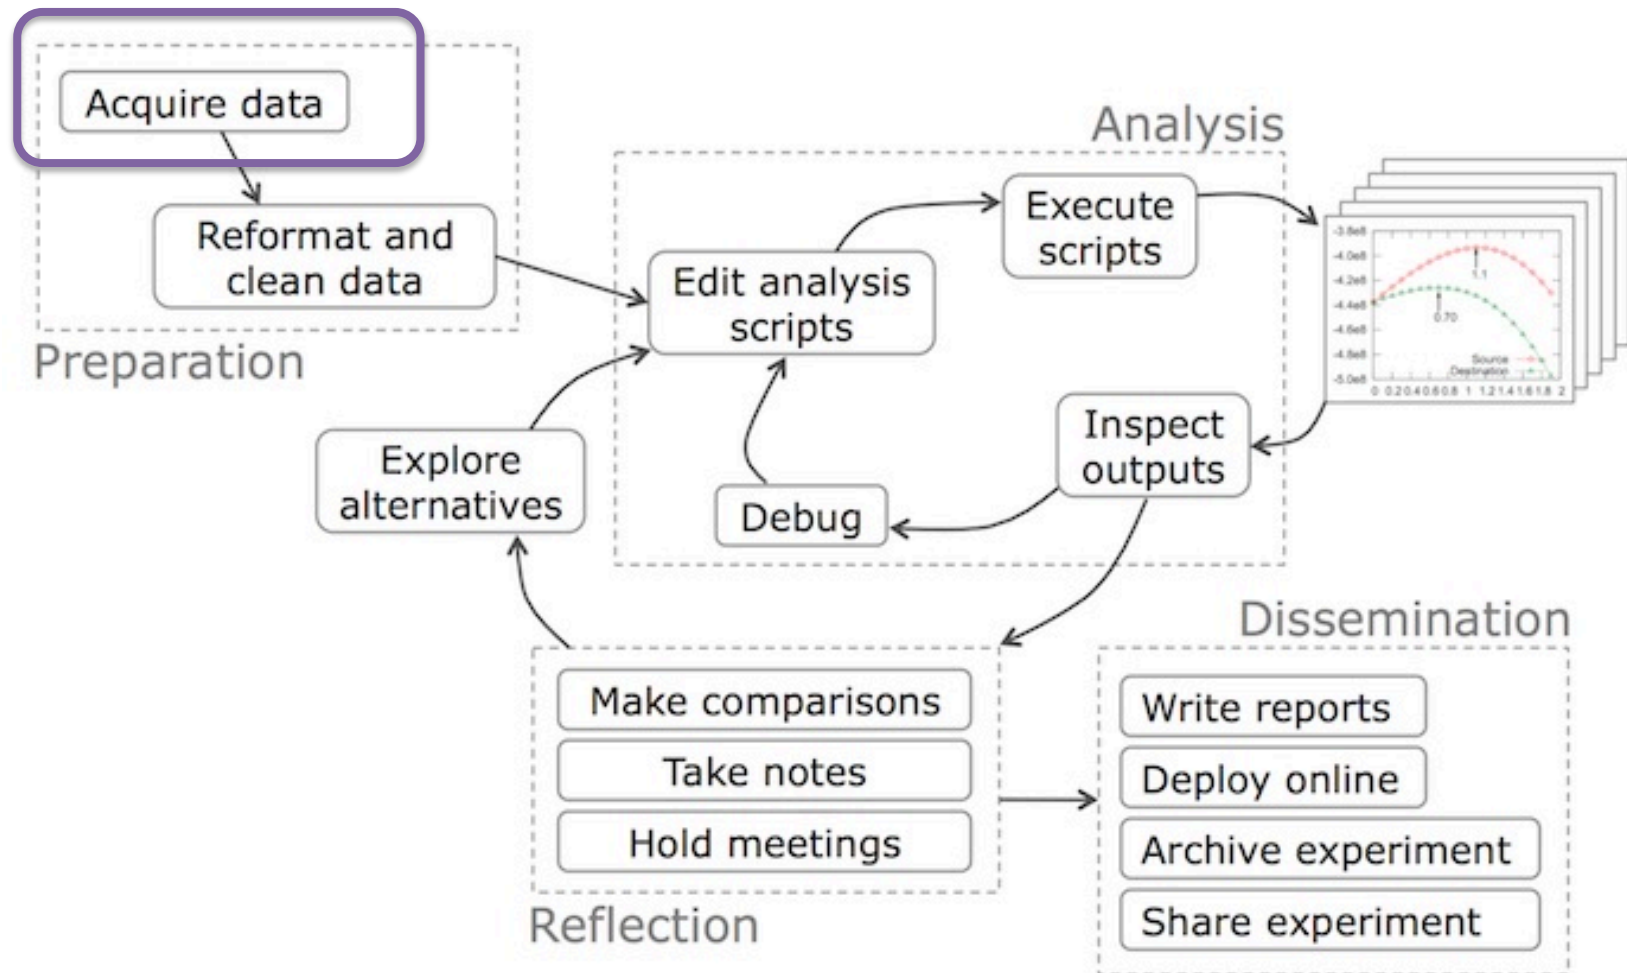


Paper Contributions

Paper's Contributions

- Identify an **relevant, open-source dataset** with subjective classification of food item
- Define an objective method to classify food items into **nutrient-driven food groups**
- Parallelize our methods using a **scalable framework** such as Apache Spark's **based on MapReduce**
- Critically **compare and contrast** the subjective and our nutrient-driven food classification

Data Workflow



Relevant Open-source Dataset

- Use a dataset with well-known and broadly used data format

NHANES: National Health and Nutrition Examination Survey

- Medical, demographic, and dietary records
- Available to the public for free
- *Contains subjective food groups provided by USDA*

Data available at: http://www.cdc.gov/nchs/nhanes/nhanes_questionnaires.htm

NHANES Dietary Data


- Dietary intake of 64,653 Americans
- 7,494 unique food items
- 1,587,750 food entries
- 46 nutrient features for each food item
 - Macronutrients (e.g., fats, carbohydrates)
 - Micronutrients (e.g., vitamins, minerals)

NHANES Dietary Data

- Dietary intake of 64,653 Americans
- 7,494 unique food items
- 1,587,750 food entries
- 46 nutrient features for each food item
 - Macronutrients (e.g., fats, carbohydrates)
 - Micronutrients (e.g., vitamins, minerals)

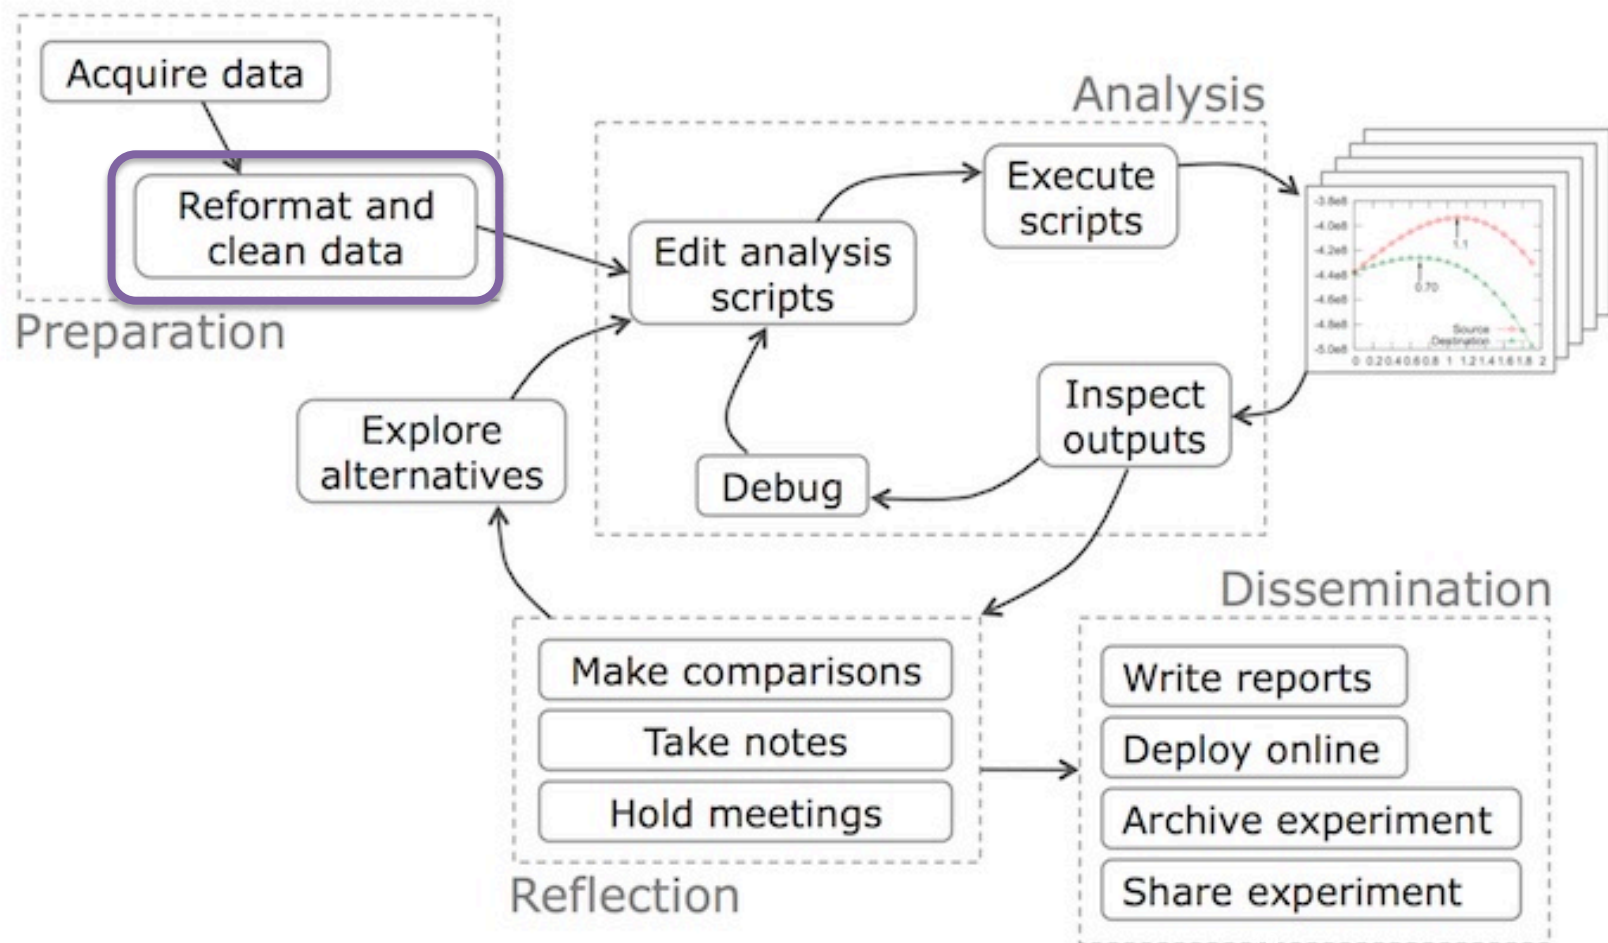
Structure of Dietary Data Item

<143672, 92510000, 3, 0, 8:15am, 7, 3, 10.1, 4, 3.45, 10, 178, ...>



- Participant ID
- Meta Data
- Macronutrients
- **USDA Food Code**
- Micronutrients

Data Workflow



NHANES Snapshot

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	<i>0</i>	244	8	107
<i>Milk</i>	<i>0</i>	122	4	54
<i>Milk</i>	<i>0</i>	100	3	44
<i>Milk</i>	<i>0</i>	300	10	132
<i>Milk</i>	<i>0</i>	10	0	4
<i>Milk</i>	<i>0</i>	93	3	41
<i>Cereal</i>	<i>0</i>	30	2	231
<i>Cereal</i>	<i>10</i>	60	6	462
<i>Cereal</i>	<i>0</i>	0	3	300
<i>Cereal</i>	<i>0</i>	10		77
<i>Steak</i>	<i>0</i>	256	62	146

Missing Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	244	8	107
<i>Milk</i>	0	122	4	54
<i>Milk</i>	0	100	3	44
<i>Milk</i>	0	300	10	132
<i>Milk</i>	0	10	0	4
<i>Milk</i>	0	93	3	41
<i>Cereal</i>	0	30	2	231
<i>Cereal</i>	10	60	6	462
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	256	62	146

Weight-Based Nutrient Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	244	8	107
<i>Milk</i>	0	122	4	54
<i>Milk</i>	0	100	3	44
<i>Milk</i>	0	300	10	132
<i>Milk</i>	0	10	0	4
<i>Milk</i>	0	93	3	41
<i>Cereal</i>	0	30	2	231
<i>Cereal</i>	10	60	6	462
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	256	62	146

Redundant Data

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
Milk	0	244	8	107
Milk	0	122	4	54
Milk	0	100	3	44
Milk	0	300	10	132
Milk	0	10	0	4
Milk	0	93	3	41
Cereal	0	30	2	231
Cereal	10	60	6	462
Cereal	0	0	3	300
Cereal	0	10		77
Steak	0	256	62	146

Features on Different Scales

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	244	8	107
<i>Milk</i>	0	122	4	54
<i>Milk</i>	0	100	3	44
<i>Milk</i>	0	300	10	132
<i>Milk</i>	0	10	0	4
<i>Milk</i>	0	93	3	41
<i>Cereal</i>	0	30	2	231
<i>Cereal</i>	10	60	6	462
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	256	62	146

Preprocessing: Missing Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Cereal</i>	0	10		77

- Standard Approach:
 - Fill in missing values e.g., median or zero
- Drawback:
 - Introduces artificial data which creates bias in the data
- Our Solution:
 - Redundant food entries allows us to discard entries with missing values
- Observations:
 - 9,586 of 1,587,750 entries removed
 - 0 unique foods lost

25

Preprocessing: Missing Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	244	8	107
<i>Milk</i>	0	122	4	54
<i>Milk</i>	0	100	3	44
<i>Milk</i>	0	300	10	132
<i>Milk</i>	0	10	0	4
<i>Milk</i>	0	93	3	41
<i>Cereal</i>	0	30	2	231
<i>Cereal</i>	10	60	6	462
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	256	62	146

Preprocessing: Weight-Based Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
Milk	0	244	8	107
Milk	0	122	4	54
Milk	0	100	3	44
Milk	0	300	10	132
Milk	0	10	0	4
Milk	0	93	3	41

- Nutrient values are based on the weight of the food entry.
 - Unfair comparison between foods
 - Food entries of the same type do not match

Preprocessing: Weight-Based Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Cereal</i>	0	0	3	300

- Standard Approach:
 - Normalize with respect to weight – divide entry by weight
- Drawback:
 - Some entries have a weight of “0”
- Our Solution:
 - Redundant food entries allows us to discard entries with a weight of “0”
- Observations:
 - 10,983 of 1,587,750 entries removed
 - 0 unique foods lost

Preprocessing: Weight-Based Values

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	1 (244)	0.033	0.439
<i>Milk</i>	0	1 (122)	0.033	0.443
<i>Milk</i>	0	1 (100)	0.030	0.440
<i>Milk</i>	0	1 (300)	0.033	0.440
<i>Milk</i>	0	1 (10)	0	0.400
<i>Milk</i>	0	1 (93)	0.032	0.441
<i>Cereal</i>	0	1 (30)	0.067	7.70
<i>Cereal</i>	10	1 (60)	0.100	7.70
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	1 (256)	0.242	0.570

Preprocessing: Redundant Data

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
Milk	0	1 (244)	0.033	0.439
Milk	0	1 (10)	0	0.400
Cereal	0	1 (30)	0.067	7.70
Cereal	10	1 (60)	0.100	7.70

- Standard Approach:
 - Select one entry for each unique food item
- Drawback:
 - Nutrient densities don't always match for the same food item – How do we pick a representative entry?
 - Modification codes
 - Rounding error

Preprocessing: Redundant Data

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
Milk	0	1 (244)	0.033	0.439
Milk	0	1 (10)	0	0.400
Cereal	0	1 (30)	0.067	7.70
Cereal	10	1 (60)	0.100	7.70

- Our Solution:
 - Remove entries with a non-zero modification code
 - Average top 5 entries for each food when sorted by weight
- Observations:
 - 32 unique foods lost

Preprocessing: Redundant Data

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	1 (244)	0.033	0.439
<i>Milk</i>	0	1 (122)	0.033	0.443
<i>Milk</i>	0	1 (100)	0.030	0.440
<i>Milk</i>	0	1 (300)	0.033	0.440
<i>Milk</i>	0	1 (10)	0	0.400
<i>Milk</i>	0	1 (93)	0.032	0.441
<i>Cereal</i>	0	1 (30)	0.067	7.70
<i>Cereal</i>	10	1 (60)	0.100	7.70
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	1 (256)	0.242	0.570

Preprocessing: Redundant Data

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	1 (244)	0.032	0.441
<i>Milk</i>	0	1 (122)	0.033	0.443
<i>Milk</i>	0	1 (100)	0.030	0.440
<i>Milk</i>	0	1 (300)	0.033	0.440
<i>Milk</i>	0	1 (10)	0	0.400
<i>Milk</i>	0	1 (93)	0.032	0.441
<i>Cereal</i>	0	1 (30)	0.067	7.70
<i>Cereal</i>	10	1 (60)	0.100	7.70
<i>Cereal</i>	0	0	3	300
<i>Cereal</i>	0	10		77
<i>Steak</i>	0	1 (256)	0.242	0.570

Preprocessing: Different Scales

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	0	1	0.032	0.441
<i>Cereal</i>	0	1	0.067	7.70
<i>Steak</i>	0	1	0.242	0.570

- Standard Approach:
 - Standardize – divide by largest value
- Drawback:
 - Removes effects of highly skewed feature distributions

Preprocessing: Different Scales

Food	Mod Code	Weight (g)	Protein (sd)	Sodium (sd)
<i>Milk</i>	0	1	-0.726	-0.593
<i>Cereal</i>	0	1	-0.416	1.16
<i>Steak</i>	0	1	1.13	-0.561

- Our Solution:
 - **Z-score standardization:** convert original values into values indicating how many standard deviations above or below a value is from the mean of the original distribution
- Observations:
 - Feature scales are similar, but not the same
 - Distributions are not dramatically altered



Preprocessing: Different Scales

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	<i>0</i>	1	0.032	0.441
<i>Cereal</i>	<i>0</i>	1	0.067	7.70
<i>Steak</i>	<i>0</i>	1	0.242	0.570

Food	Mod Code	Weight (g)	Protein (sd)	Sodium (sd)
<i>Milk</i>	<i>0</i>	1	-0.726	-0.593
<i>Cereal</i>	<i>0</i>	1	-0.416	1.16
<i>Steak</i>	<i>0</i>	1	1.13	-0.561

Example of Standardization

- Macro-nutrient, protein, with its content distribution before and after standardization



NHANSE Snapshot After Preparation

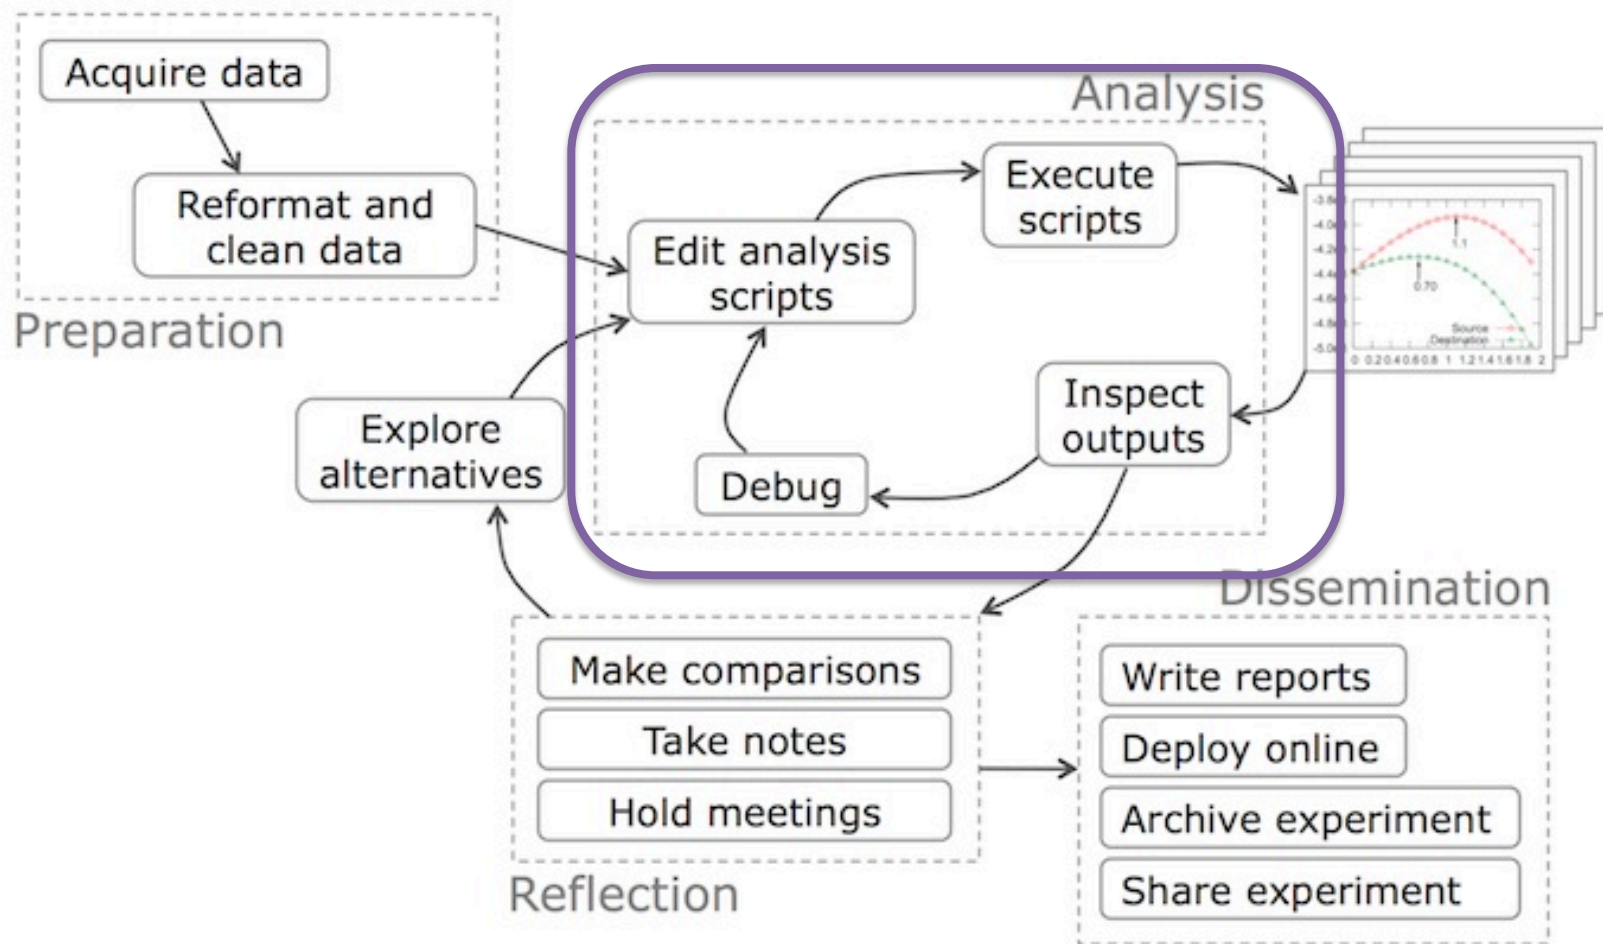
Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
<i>Milk</i>	<i>0</i>	<i>1</i>	<i>-0.726</i>	<i>-0.593</i>
<i>Milk</i>	<i>0</i>	<i>1</i>	<i>0.033</i>	<i>0.443</i>
<i>Milk</i>	<i>0</i>	<i>1</i>	<i>0.030</i>	<i>0.440</i>
<i>Milk</i>	<i>0</i>	<i>1</i>	<i>0.033</i>	<i>0.440</i>
<i>Milk</i>	<i>0</i>	<i>1</i>	<i>0</i>	<i>0.400</i>
<i>Milk</i>	<i>0</i>	<i>1</i>	<i>0.032</i>	<i>0.441</i>
<i>Cereal</i>	<i>0</i>	<i>1</i>	<i>-0.416</i>	<i>1.16</i>
<i>Cereal</i>	<i>10</i>	<i>1</i>	<i>0.100</i>	<i>7.70</i>
<i>Cereal</i>	<i>0</i>	<i>0</i>	<i>3</i>	<i>300</i>
<i>Cereal</i>	<i>0</i>	<i>10</i>		<i>77</i>
<i>Steak</i>	<i>0</i>	<i>1</i>	<i>1.13</i>	<i>-0.561</i>

NHANSE Snapshot After Preparation

Food	Mod Code	Weight (g)	Protein (g)	Sodium (mg)
Milk	0	1	-0.726	-0.593
Milk	0	1	0.033	0.443
Milk	0	1	0.033	0.440
Milk	0	1	0.033	0.440
Milk	0	1	0.033	0.400
Milk	0	1	0.032	0.441
Cereal	0	1	-0.416	1.16
Cereal	10	1	0.100	7.70
Cereal	0	0	3	300
Cereal	0	10		77
Steak	0	1	1.13	-0.561

BEFORE: 7,494 unique food items
AFTER: 7,462 unique food items

Data Workflow




Selecting a Good Clustering Algorithm

K-Means

Hierarchical

DBSCAN




Feature	K-Means	Hierarchical	DBSCAN
<i>Resource Efficient</i>	Yes	No	Yes
<i>Noise Insensitive</i>	No	No	Yes
<i>Outlier Detection</i>	No	No	Yes
<i>Spheroid Clusters</i>	Yes	Yes	Yes
<i>Non-Spheroid Clusters</i>	No	Yes	Yes
<i>Undefined Cluster Count</i>	No	Yes	Yes

Image: <http://scikit-learn.org/stable/modules/clustering.html>

DBSCAN Algorithm


Epsilon: Distance around point for counting neighbors




Min_pts: Minimum neighbors for “core points”

Together, these parameters define the minimum density for a cluster

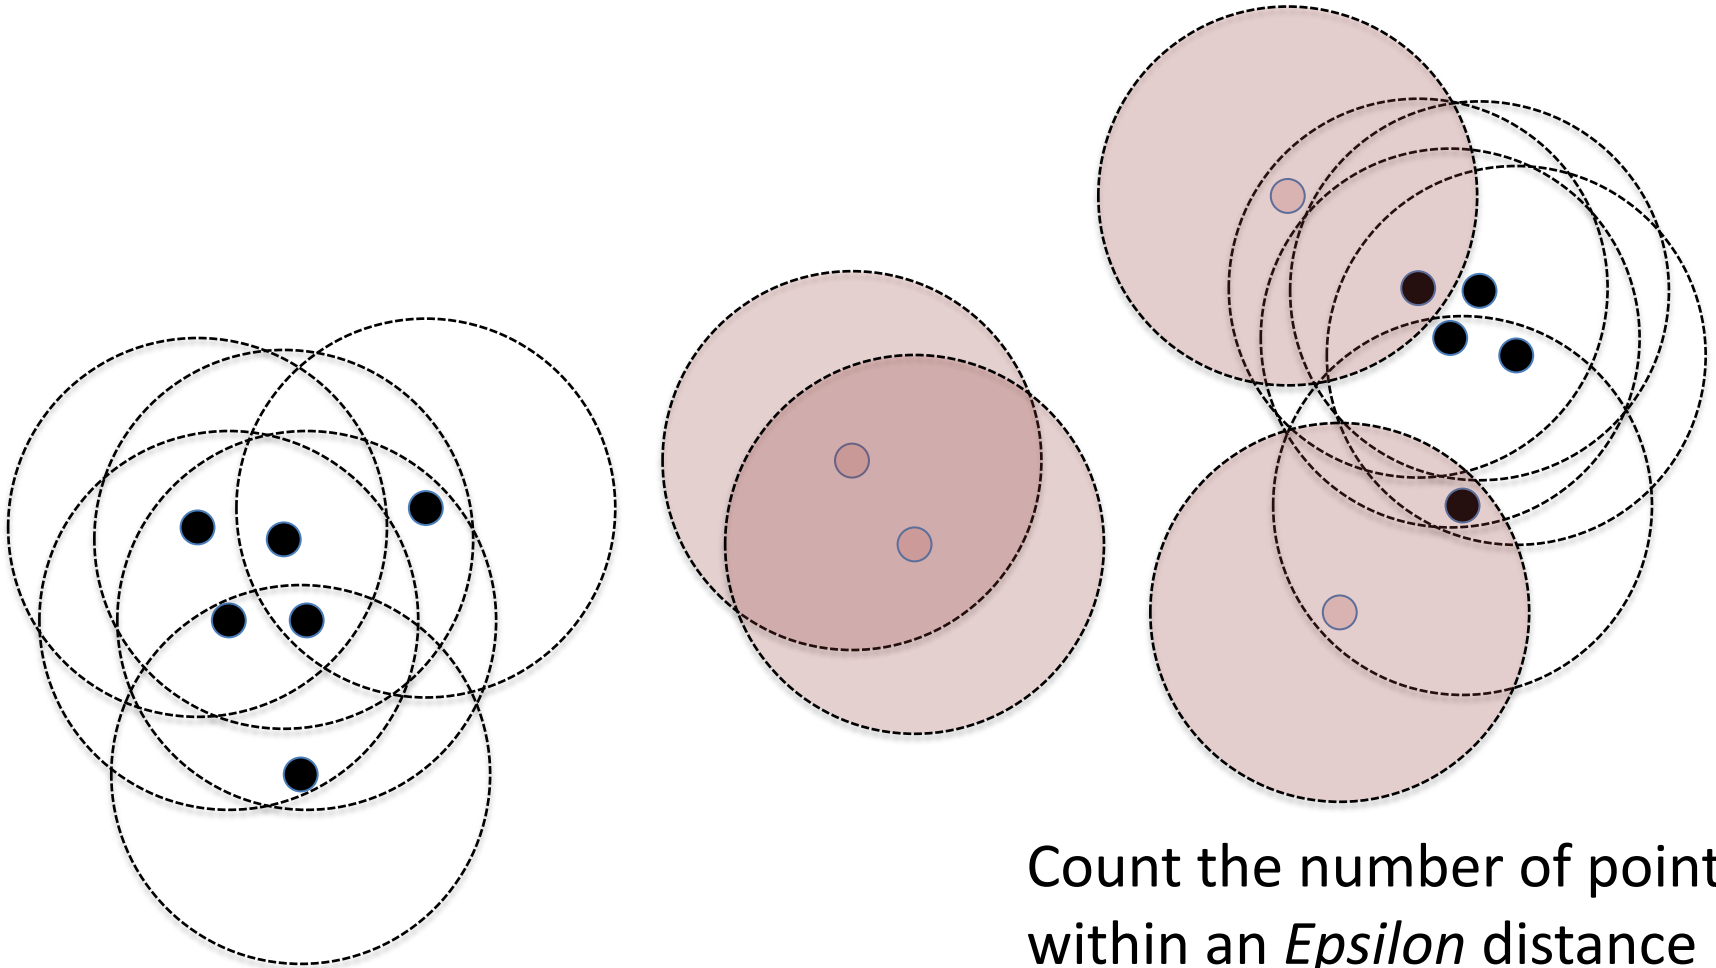
DBSCAN Algorithm



DBSCAN Algorithm: Border Points



DBSCAN Algorithm: Noise Points




DBSCAN Algorithm



Identify **Core**, **Border**, and **Noise** points

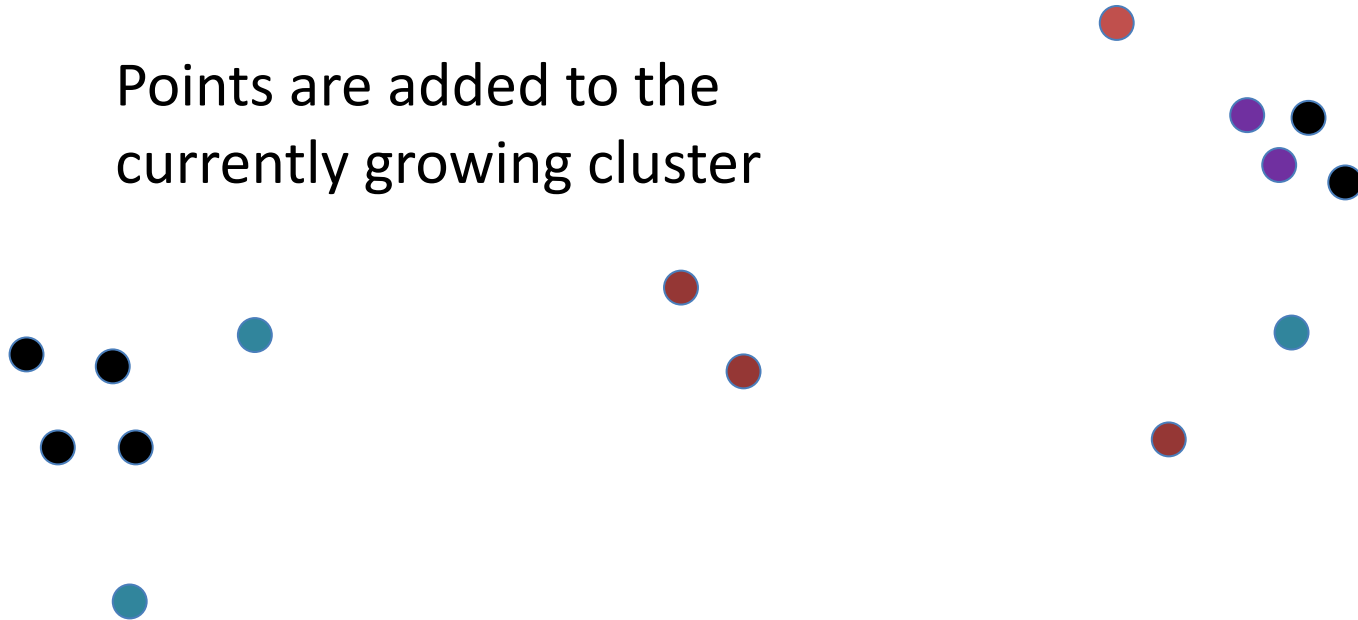
DBSCAN Algorithm



Start with a **Core** point and begin growing the clusters.


DBSCAN Algorithm

Points are added to the currently growing cluster




DBSCAN Algorithm

When no more points are reachable, the next un-clustered **Core** point is chosen





DBSCAN Algorithm



This pattern continues until only **Noise** points remain un-clustered.

DBSCAN Algorithm




BDSCAN Demo

- Visualizing a DBSCAN clustering
 - Interactive interface at:
 - <https://www.naftaliharris.com/blog/visualizing-dbscan-clustering/>


Bad Clustering

- Very Wide
- Not well separated



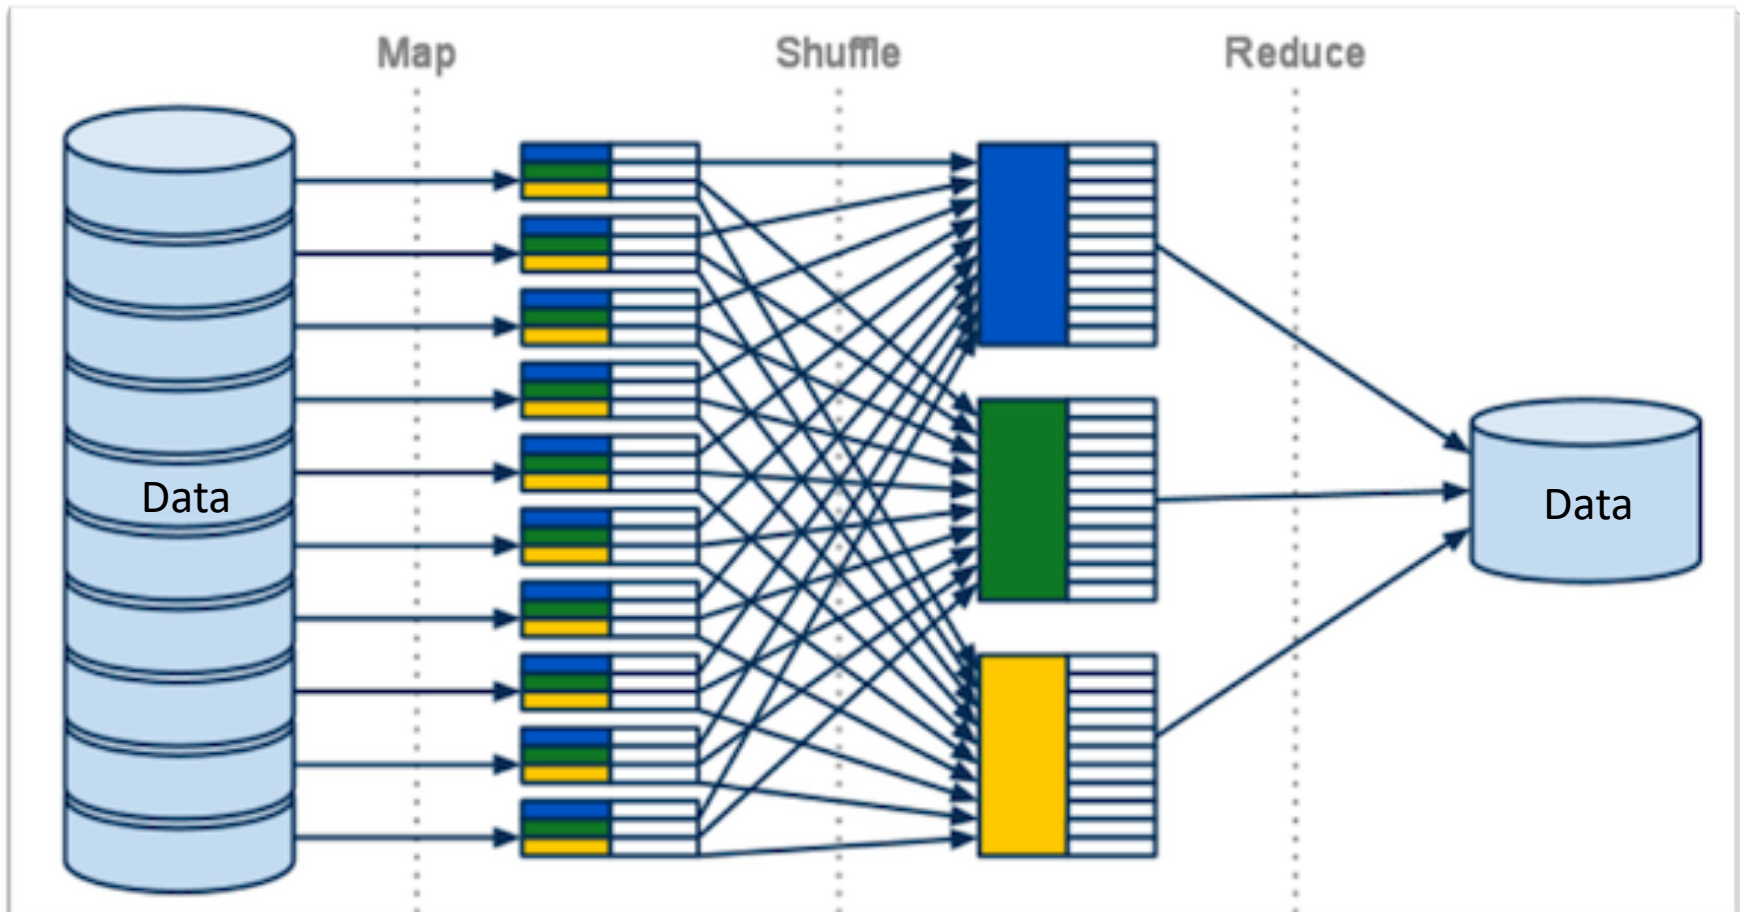
Good Clustering

- Dense
- Well separated

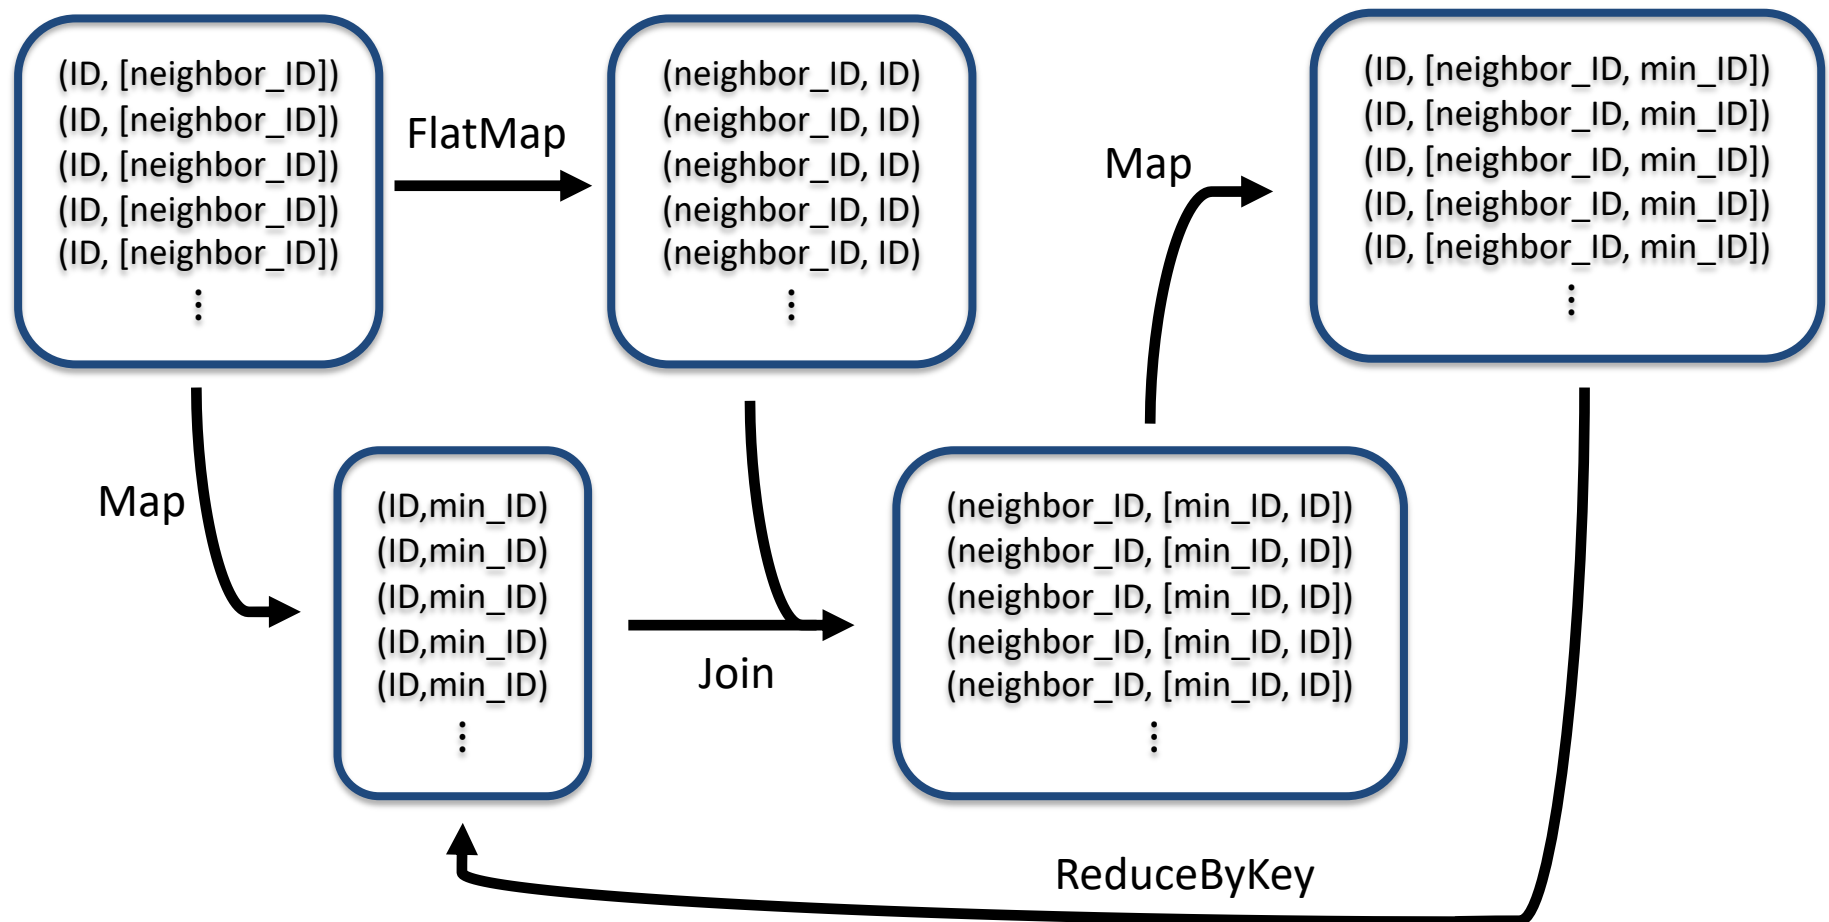


MapReduce Paradigm

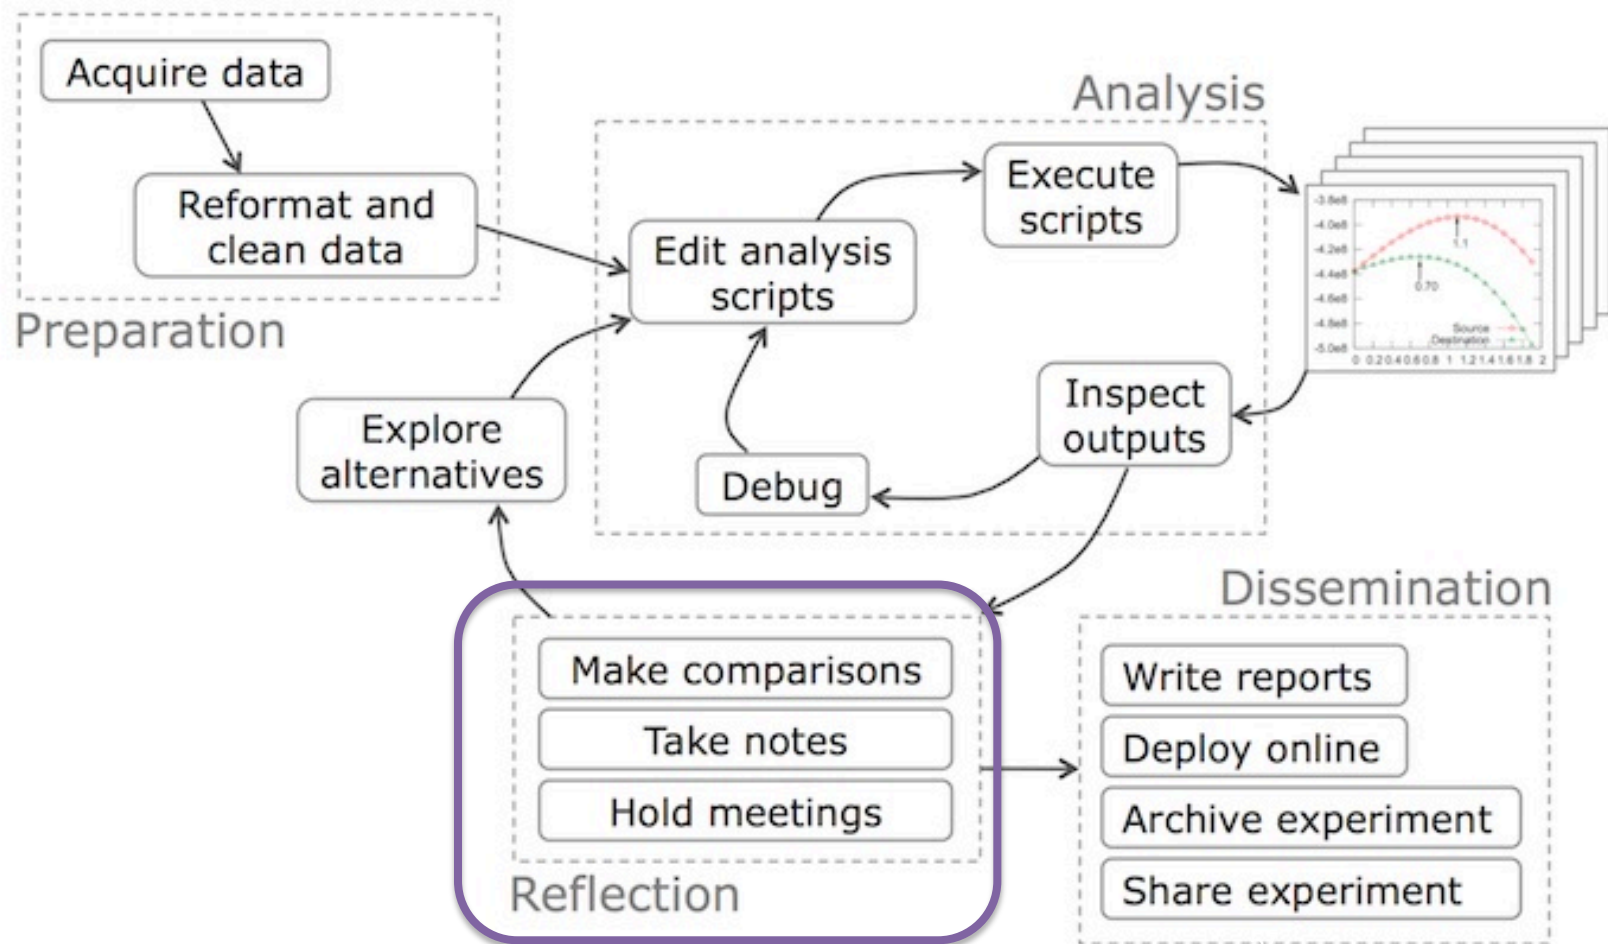
Image: <https://aimotion.blogspot.com/2012/08/introduction-to-recommendations-with.html>



Parallel DBSCAN with Spark



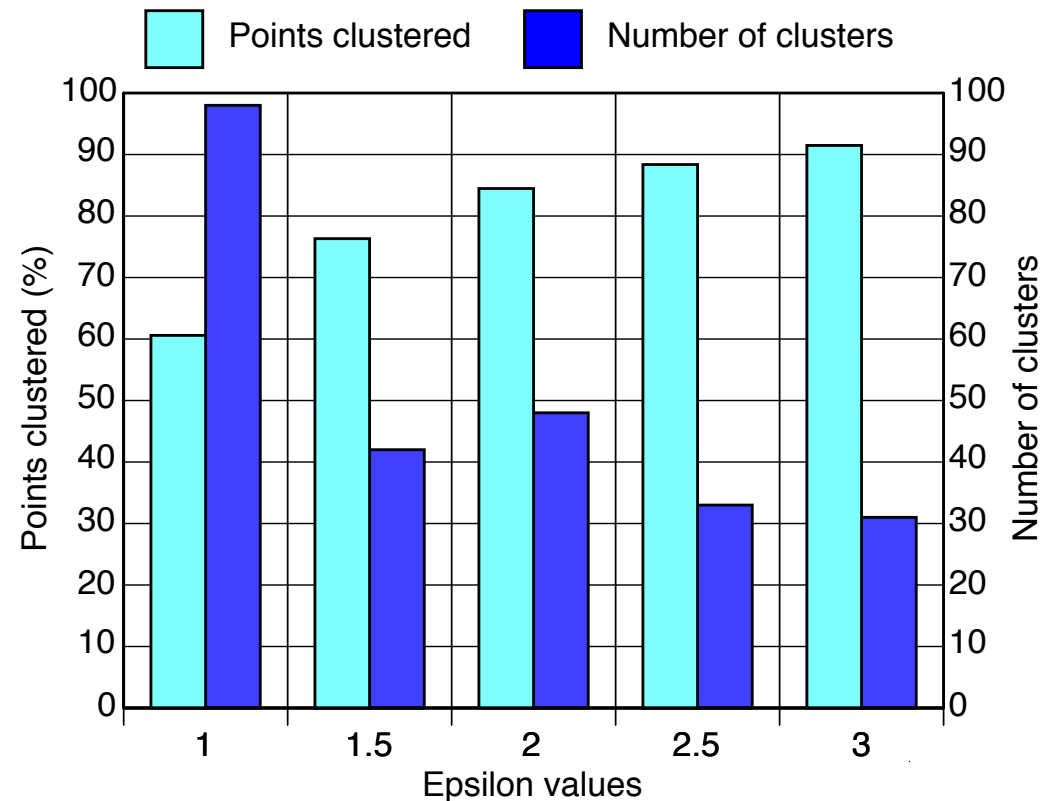
Data Workflow



Courtesy of Philip Guo: goo.gl/y42rp1


Experiment Setting and Metrics

- DBSCAN settings:
 - Epsilon = 1.0
 - Min_pts = 4
 - Euclidean distance



Experiment Setting and Metrics

- DBSCAN settings:
 - Epsilon = 1.0
 - Min_pts = 4
 - Euclidean distance
- Metrics of success:
 - Cluster diameter (min)
 - Cluster separation (max)




Experiment Setting and Metrics

- DBSCAN settings:
 - Epsilon = 1.0
 - Min_pts = 4
 - Euclidean distance
- Metrics of success:
 - Cluster diameter (min)
 - Cluster separation (max)
- Comparisons:
 - **Our clustering**
 - **USDA code clustering with 2-digit code**
 - **USDA code clustering with 3-digit code**


2-digit
92510000
3-digit

A diagram showing the number 92510000. A bracket above the first two digits (92) is labeled '2-digit'. A bracket below the first three digits (925) is labeled '3-digit'.

Clustering Results

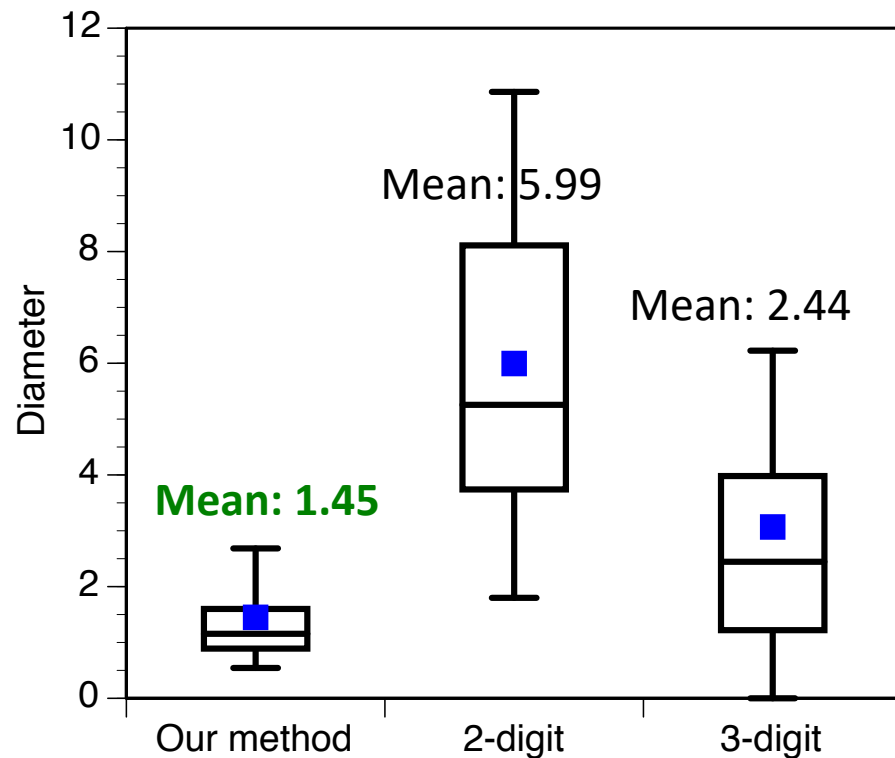


Clustering Results




Clustering Results

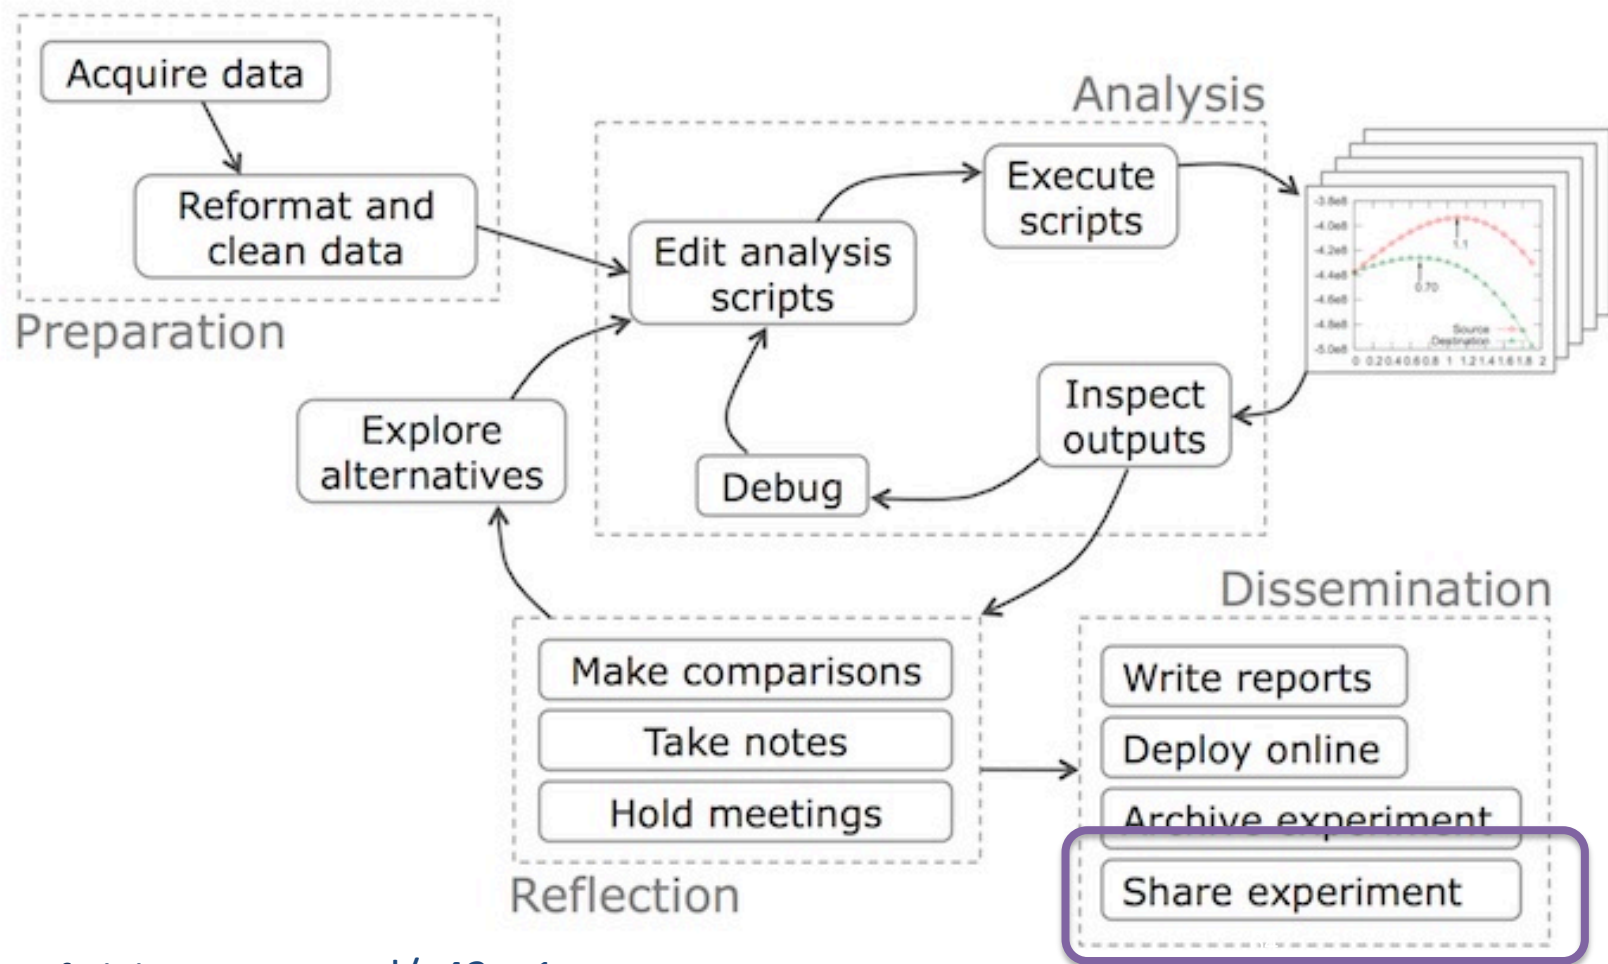
Cluster Diameter → Lower=Better



Cluster Separation → Higher=Better



Data Workflow



Courtesy of Philip Guo: goo.gl/y42rp1

Lessons Learned

- Using the traditional USDA classification of food items is misleading and can poorly advise patients with health issues
“Eat less fat!” “but USDA codes aren’t based on nutrient content!”
- We propose a comprehensive data analysis workflow for NHANES dietary data
- Our approach clusters food items based **EXCLUSIVELY** on their nutritional content (i.e., micro- and macro-nutrients)
- Our methods is **scalable** (based on MapReduce) and produces **denser** and **better separated** food groups than USDA
 - Denser cluster diameter: 1.45 vs. 5.99 and 3.07
 - Better separated cluster separation: 1.83 vs. 0.34 and 0.39
- Our approach can provide a better indication of food group quality when advising patients with dietary restrictions